

Week 7 Exercises

Exercise 1 First draft

Make sure to upload your “first draft” to Canvas by Sunday night. Also email the PDF to your reviewer below.

Author	Reviewer
Caroline	David
Sarah	Taylor
Taylor	Shaka
Bryson	Sarah
David	Bryson
Shaka	Caroline

Exercise 2 Provide feedback

Provide comments on your assigned paper. When providing feedback on colleagues’ work, I like the RAP model [here](#). Most importantly, in this model, “focus time and effort on the *highest points of leverage* for improving the document.”

The goal is to make suggestions that improve the final product at the end of the semester.

In addition to other comments, please make sure to provide feedback on these three items:

1. Does the paper make a clear *descriptive* claim? It might be helpful to review the key readings on descriptive research [here](#), such as Gerring (2012), Holmes et al. (2024), and de Kadt and Grzymala-Busse (2025).
2. Does the paper make a compelling empirical argument for that descriptive claim?
3. Does the paper clearly explain the substantive/theoretical/normative importance of that descriptive claim? It’s important that the defense not rely on any assumed/implied *causal* relationship among the variables.

Exercise 3 Eight schools

In my slides, I motivated hierarchical models with the “309 problem.” There’s a similar motivational problem called the “eight schools” problem. This example is worth studying carefully.

1. Read this blog post (by Phil Price) on Andrew Gelman’s blog: “[Everything I need to know about Bayesian statistics, I learned in eight schools.](#)”
2. Read the pp. 119-124 of *BDA3* on the eight schools problem.
3. Run the R code [here](#). Note that the notation in my Stan model differs from BDA3.
4. In the model, σ captures the *similarity* of the treatment effects across school districts. I give this parameter a [half-Cauchy prior](#) with a scale of 20. (See the note below.) How does changing the scale of the half-Cauchy prior in the Stan model affect the estimates of σ ?
 - a. Try setting the scale of the half-Cauchy to 5 instead of 20.
 - b. Try setting it too 100. Hint: you can use `sigma_draws <- spread_draws(fit, sigma)` to get a convenient data frame of posterior simulations of σ .

Make sure you understand the basic eight schools example and can discuss it in parallel to the 309 example in the slides.

So what’s a half-Cauchy?

A Cauchy distribution is a location-scale t distribution with one degree of freedom. The *half-Cauchy* is centered at zero and *folded over*, so that it is just the rightward half. It has only a scale parameter that corresponds to something like the “expected” value of σ .¹ The expected value of the half-Cauchy does not exist, but the median equals the scale parameter. This is a useful *weakly informative* prior for variance parameters that capture similarity in hierarchical models.

In this case, the student-level SAT scores range from 200 to 800, and the treatment effects are somewhere around a 10-point improvement. In the Stan model, σ is the SD around that 10-point effect across schools. $\sigma = 20$ would be huge variation; $\sigma = 3$ or $\sigma = 5$ seem “about right.” So I think of a half-Cauchy with a scale of 20 as a “weakly informative” prior that only rules out absurdly large values of σ and gently encourages smaller rather than larger values.

¹The scale parameters of the Cauchy is also denoted as σ , so I just call it the “scale” here.

